

Summary Dialogue on Memory Virtualization

Student: *(Gulps)* Wow, that was a lot of material.

Professor: Yes, and?

Student: Well, how am I supposed to remember it all? You know, for the exam?

Professor: Goodness, I hope that's **not why you are trying to remember it.**

Student: Why should I then?

Professor: Come on, I thought you knew better. You're trying to learn something here, so that when you **go off into the world**, you'll understand how systems actually work.

Student: Hmm... can you give an example?

Professor: Sure! One time back in graduate school, my friends and I were measuring **how long memory accesses took**, and once in a while the numbers were way **higher than we expected**; we thought all the data was fitting nicely into the second-level hardware cache, you see, and thus should have been really fast to access.

Student: *(nods)*

Professor: We couldn't figure out what was going on. So what do you do in such a case? **Easy, ask a professor!** So we went and asked one of our professors, who looked at the graph we had produced, and **simply said "TLB"**. Aha! Of course, TLB misses! Why didn't we think of that? Having a **good model of how virtual memory works** helps diagnose all sorts of interesting performance problems.

Student: I think I see. I'm trying to build these mental models of how things work, so that when I'm out there working on my own, I won't be surprised when a system doesn't quite behave as expected. I should **even be able to anticipate how the system will work** just by thinking about it.

Professor: Exactly. So what have you learned? What's in your mental model of how virtual memory works?

Student: Well, I think I now have a pretty good idea of what happens when memory is referenced by a process, which, as you've said many times, happens on each instruction fetch as well as explicit loads and stores.

Professor: Sounds good — tell me more.

Student: Well, one thing I'll always remember is that the addresses we see in a user program, written in C for example...

Professor: What other language is there?

Student: (continuing) ... Yes, I know you like C. So do I! Anyhow, as I was saying, I now really know that all addresses that **we can observe within a program are virtual addresses**; that I, as a programmer, am just given this illusion of where data and code are in memory. I used to think it was cool that I could print the address of a pointer, but now I find it frustrating — it's just a virtual address! I can't see the real physical address where the data lives.

Professor: Nope, the OS **definitely hides that** from you. What else?

Student: Well, I think the TLB is a really key piece, providing the system with a small hardware cache of address translations. Page tables are usually quite large and hence live in big and slow memories. Without that TLB, programs would certainly run a great deal more slowly. **Seems like the TLB truly makes virtualizing memory possible**. I couldn't imagine building a system without one! And I shudder at the thought of a program with a working set that exceeds the coverage of the TLB: with all those TLB misses, it would be hard to watch.

Professor: Yes, cover the eyes of the children! Beyond the TLB, what did you learn?

Student: I also now understand that the page table is one of those data structures you need to know about; it's **just a data structure, though**, and that means almost any structure could be used. We started with simple structures, like arrays (a.k.a. linear page tables), and **advanced all the way up** to multi-level tables (which look like trees), and even crazier things like pageable page tables in **kernel virtual memory**. All to save a little space in memory!

Professor: Indeed.

Student: And here's one more important thing: I learned that the address translation structures **need to be flexible enough** to support what programmers want to do with their address spaces. Structures like the multi-level table are perfect in this sense; they only create table space when the user **needs a portion** of the address space, and thus there is **little waste**. Earlier attempts, like the simple base and bounds register, just weren't flexible enough; the structures need to match **what users expect and want** out of their virtual memory system.

Professor: That's a nice perspective. What about all of the stuff we learned about swapping to disk?

Student: Well, it's certainly fun to study, and good to know how page replacement works. Some of the basic policies are kind of obvious (like LRU, for example), but building a real virtual memory system seems more interesting, like

we saw in the VMS case study. But somehow, I found the mechanisms more interesting, and the policies less so.

Professor: *Oh, why is that?*

Student: *Well, as you said, in the end the best solution to policy problems is simple: buy more memory. But the mechanisms you need to understand to know how stuff really works. Speaking of which...*

Professor: *Yes?*

Student: *Well, my machine is running a little slowly these days... and memory certainly doesn't cost that much...*

Professor: *Oh fine, fine! Here's a few bucks. Go and get yourself some DRAM, cheapskate.*

Student: *Thanks professor! I'll never swap to disk again — or, if I do, at least I'll know what's actually going on!*